


# Finding Dense Subgraphs via Low-Rank Bilinear Optimization

**Ioannis Mitliagkas**



**UT Austin**

with: **Dimitris Papailiopoulos**


**Alex Dimakis**

**Constantine Caramanis**


# Densest k-Subgraph (DkS)

Given  
**graph** and a **parameter k**

Find  
**k vertices** containing **most edges**



# Densest k-Subgraph (DkS)



Given  
**graph** and a **parameter k**

Find  
**k vertices** containing **most edges**

Applications

**Community Mining**

*communities = large dense components*


**Link Spam Detection**

*dense parts of web: **spam***

**Computational biology**

*complex patterns in gene annotation graphs*


# Densest k-Subgraph (DkS)



There is a  
**5-subgraph with 10 edges**

**Q: Can you find it?**

# Densest k-Subgraph (DkS)




Given  
**graph** and a **parameter k**

Find  
**k vertices** containing **most edges**

**NP-hard**

**Hard to approximate**

# Densest k-Subgraph (DkS)



Given  
**graph** and a **parameter k**

Find  
**k vertices** containing **most edges**

**NP-hard**

**Hard to approximate**

[Khot, 2004]

\*Except in specific cases: [Arora et al 95]  
(1+ $\epsilon$ ) approx. for linear subgraphs of dense graphs



# Worst-Case Analysis

# Worst-Case Analysis

$$\text{density} = \frac{2 \cdot \# \text{ edges in subgraph}}{k} \quad (\text{av.degree})$$

$$\text{Approx} \geq \frac{\text{OPT}}{\rho}$$



# Worst-Case Analysis

$$\text{density} = \frac{2 \cdot \# \text{ edges in subgraph}}{k} \quad (\text{av.degree})$$

$$\text{Approx} \geq \frac{\text{OPT}}{\rho}$$

**After long effort,** [Feige, 2001], [Bhaskara et al., STOC '10]

**Best known ratio**

$$\text{Approx} \geq \frac{\text{OPT}}{n^{0.25}}$$

**10-factor** approx. for graphs with **10K nodes**

**100-factor** approx. for graphs with **100 Million nodes**



Known DkS guarantees are not useful in practice...  
*under worst case analysis*

Known DkS guarantees are not useful in practice...  
*under worst case analysis*

**Q1:** Provable, graph-dependent bounds?

**Q2:** DkS on billion-scale graphs?

# Beyond the Worst Case

## New DkS algorithm:

Graph-dependent bounds

In practice: **Approx**  $\geq 0.7 \cdot \mathbf{OPT}$

## Scalable


**nearly-linear times** for many real-world graphs

## Parallelizable

implementation in **MapReduce+Python**

up to **billion-edge** graphs on **800 cores on Amazon EC2**


# Our Low-Rank Framework




DkS on a graph

- Hard to solve
- Hard to approximate

# Our Low-Rank Framework



Low rank  
approximation



DkS on a graph


- Hard to solve
- Hard to approximate

DkS on constant rank graph


- Nearly-linear time solvable (!)



# Our Low-Rank Framework



Low rank  
approximation



DkS on a graph

- Hard to solve
- Hard to approximate

DkS on constant rank graph

- Nearly-linear time solvable (!)

Low-rank DkS is related to original DkS

# Results: Theory

# Graph-dependent Guarantees

$$\text{density} = \frac{2 \cdot \# \text{ edges in subgraph}}{k} \quad (\text{av.degree})$$

## Theorems:

Algorithm computes in **time**  $O(n^{d+2}/\delta)$  a  $k$ -subgraph with **density**

$$\text{OPT}_d \geq \text{OPT} \cdot 0.5 \cdot (1 - \delta) - 2|\lambda_{d+1}|$$

# Graph-dependent Guarantees

$$\text{density} = \frac{2 \cdot \# \text{ edges in subgraph}}{k} \quad (\text{av.degree})$$

## Theorems:

Algorithm computes in **time**  $O(n^{d+2}/\delta)$  a  $k$ -subgraph with **density**

$$\text{OPT}_d \geq \text{OPT} \cdot 0.5 \cdot (1 - \delta) - 2|\lambda_{d+1}|$$

If the **largest  $d$  eigenvalues** of the adjacency are **positive**  
Our algorithm computes in **time**  $O(|E| \cdot \log n + \frac{n}{\epsilon^d})$   
a  $k$ -subgraph with **density**

$$\text{OPT}_d \geq \text{OPT} \cdot (1 - \epsilon) - 2|\lambda_{d+1}|$$

# Graph-dependent Guarantees

$$\text{density} = \frac{2 \cdot \# \text{ edges in subgraph}}{k} \quad (\text{av.degree})$$

## Theorems:

Algorithm computes in **time**  $O(n^{d+2}/\delta)$  a  $k$ -subgraph with **density**

$$\text{OPT}_d \geq \text{OPT} \cdot 0.5 \cdot (1 - \delta) - 2|\lambda_{d+1}|$$

If the **largest  $d$  eigenvalues** of the adjacency are **positive**  
Our algorithm computes in **time**  $O(|E| \cdot \log n + \frac{n}{\epsilon^d})$   
a  $k$ -subgraph with **density**


$$\text{OPT}_d \geq \text{OPT} \cdot (1 - \epsilon) - 2|\lambda_{d+1}|$$

**larger  $d$**  => **better** approximation, **slower** computation

# Performance in Practice

# com-LiveJournal graph


4M nodes, 35M edges





# com-LiveJournal graph

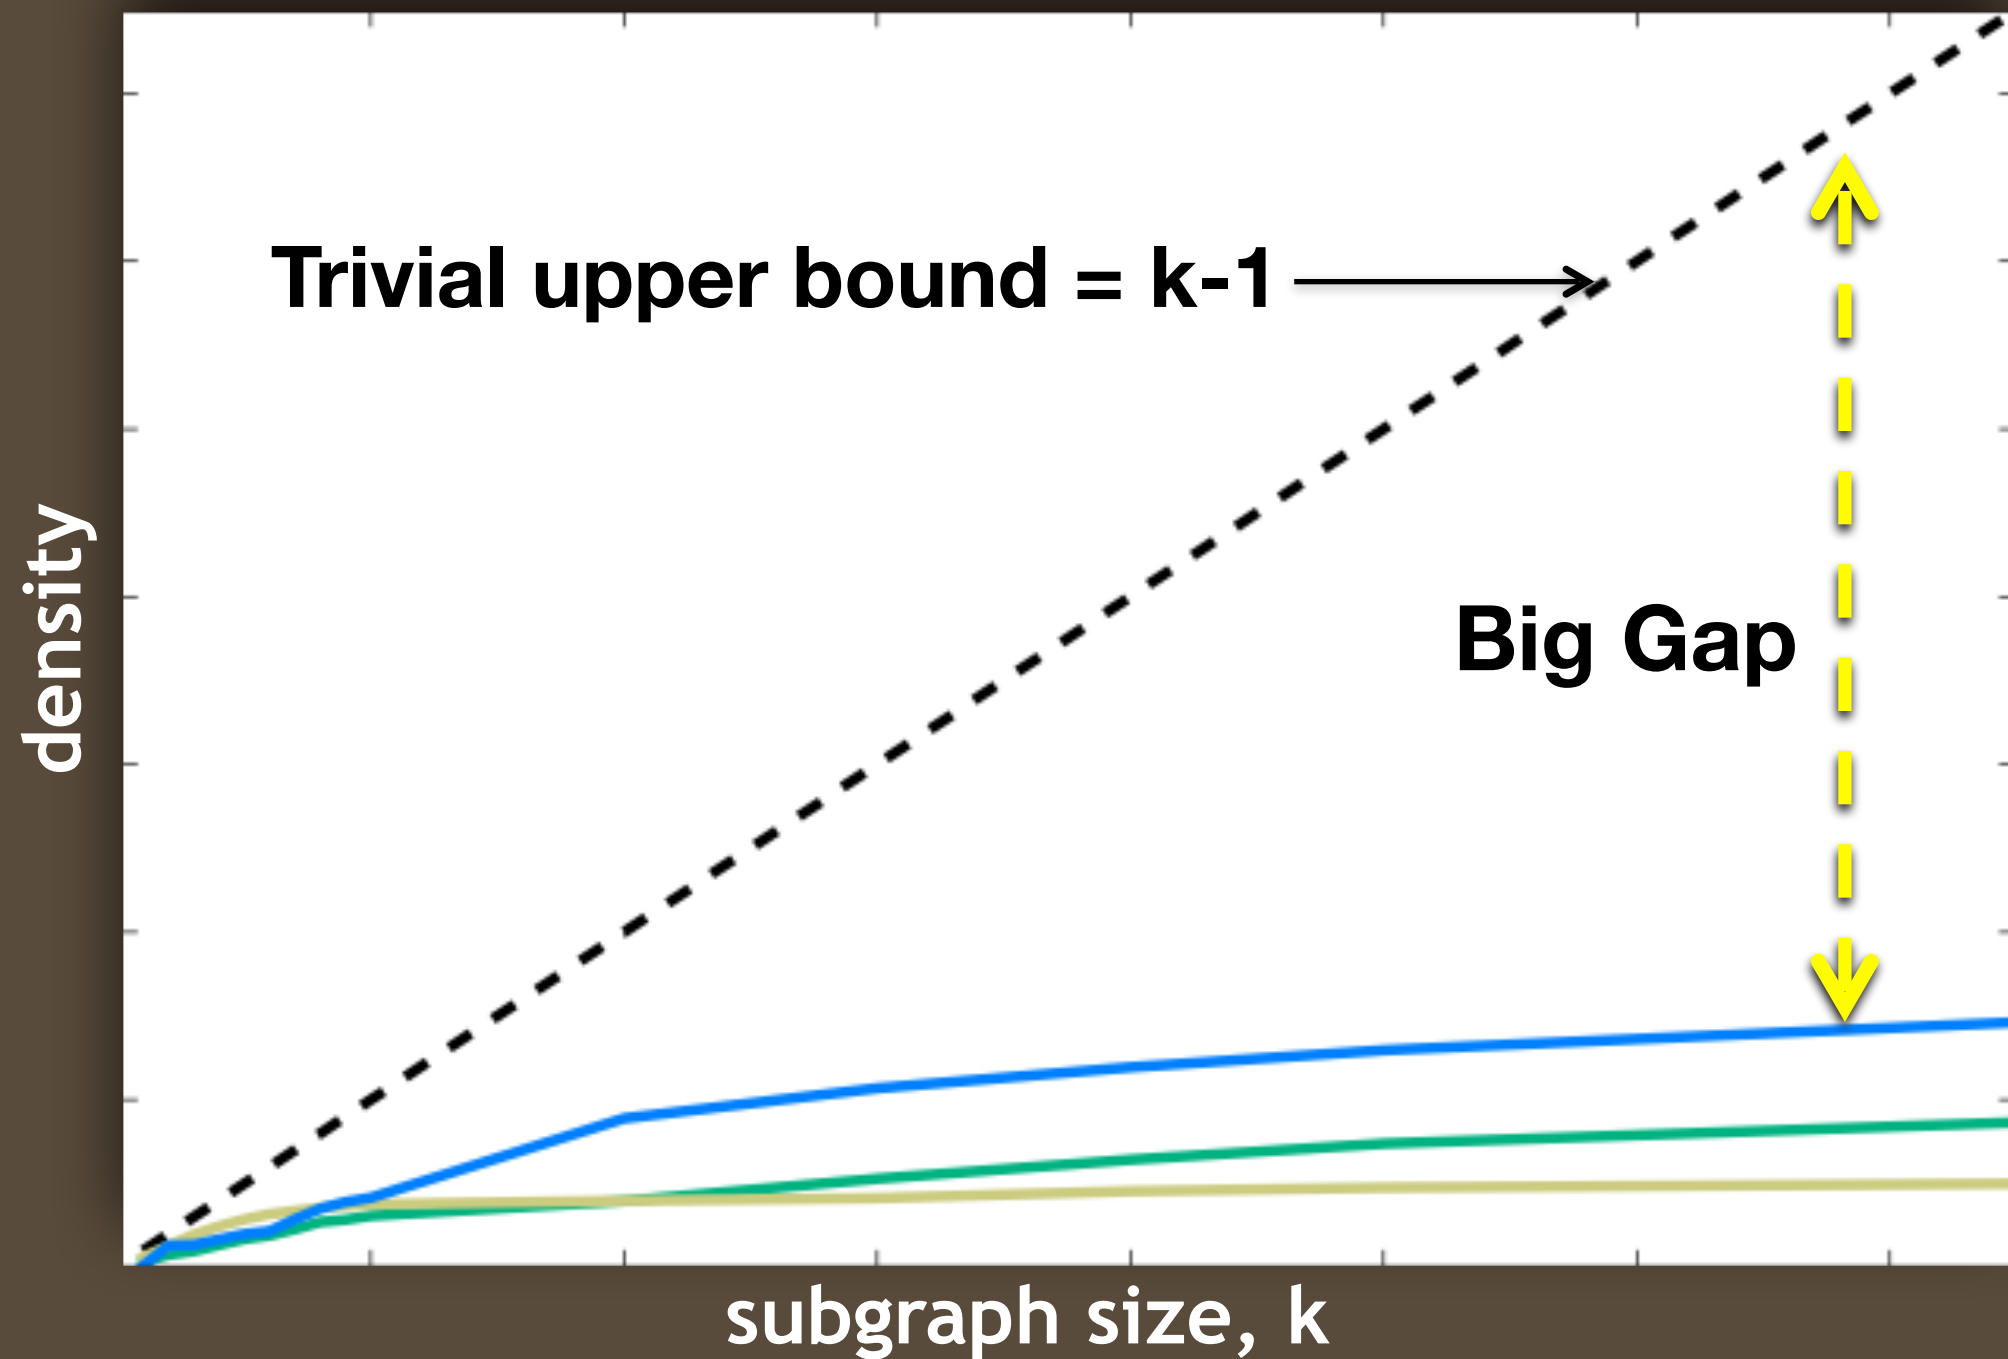
4M nodes, 35M edges



Blue: TPower JMLR'13   Green: GreedyFeige Algorithmica '01   Yellow: GreedyRavi OR'94

# com-LiveJournal graph


4M nodes, 35M edges



Blue: TPower JMLR'13 Green: GreedyFeige Algorithmica '01 Yellow: GreedyRavi OR'94

# com-LiveJournal graph

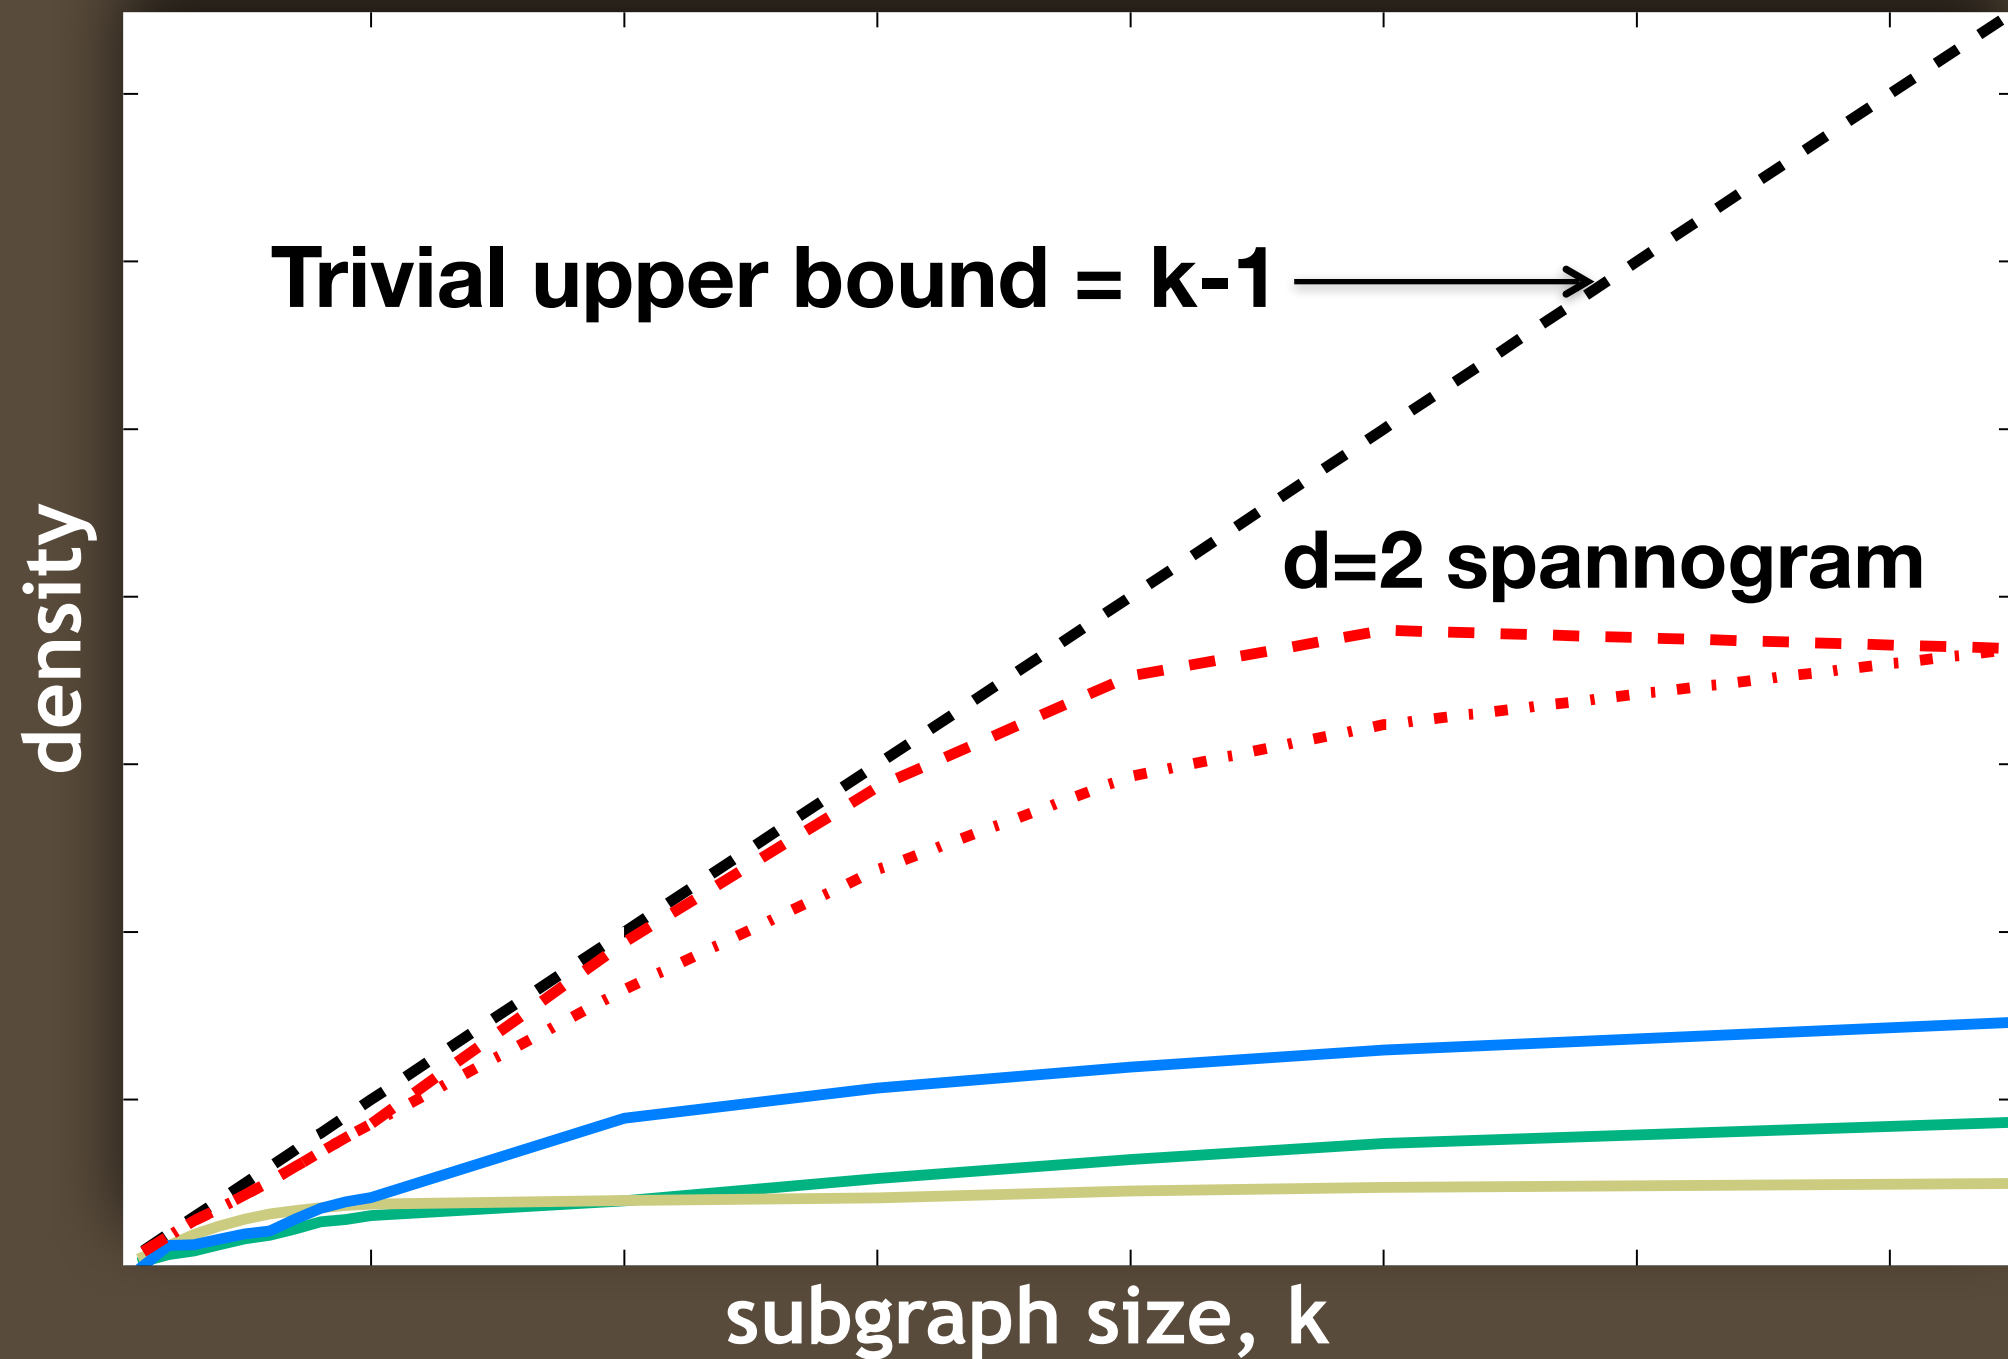
4M nodes, 35M edges



Blue: TPower JMLR'13   Green: GreedyFeige Algorithmica '01   Yellow: GreedyRavi OR'94

# com-LiveJournal graph


4M nodes, 35M edges



Blue: TPower JMLR'13   Green: GreedyFeige Algorithmica '01   Yellow: GreedyRavi OR'94

# com-LiveJournal graph


4M nodes, 35M edges



Blue: TPower JMLR'13   Green: GreedyFeige Algorithmica '01   Yellow: GreedyRavi OR'94

# com-LiveJournal graph

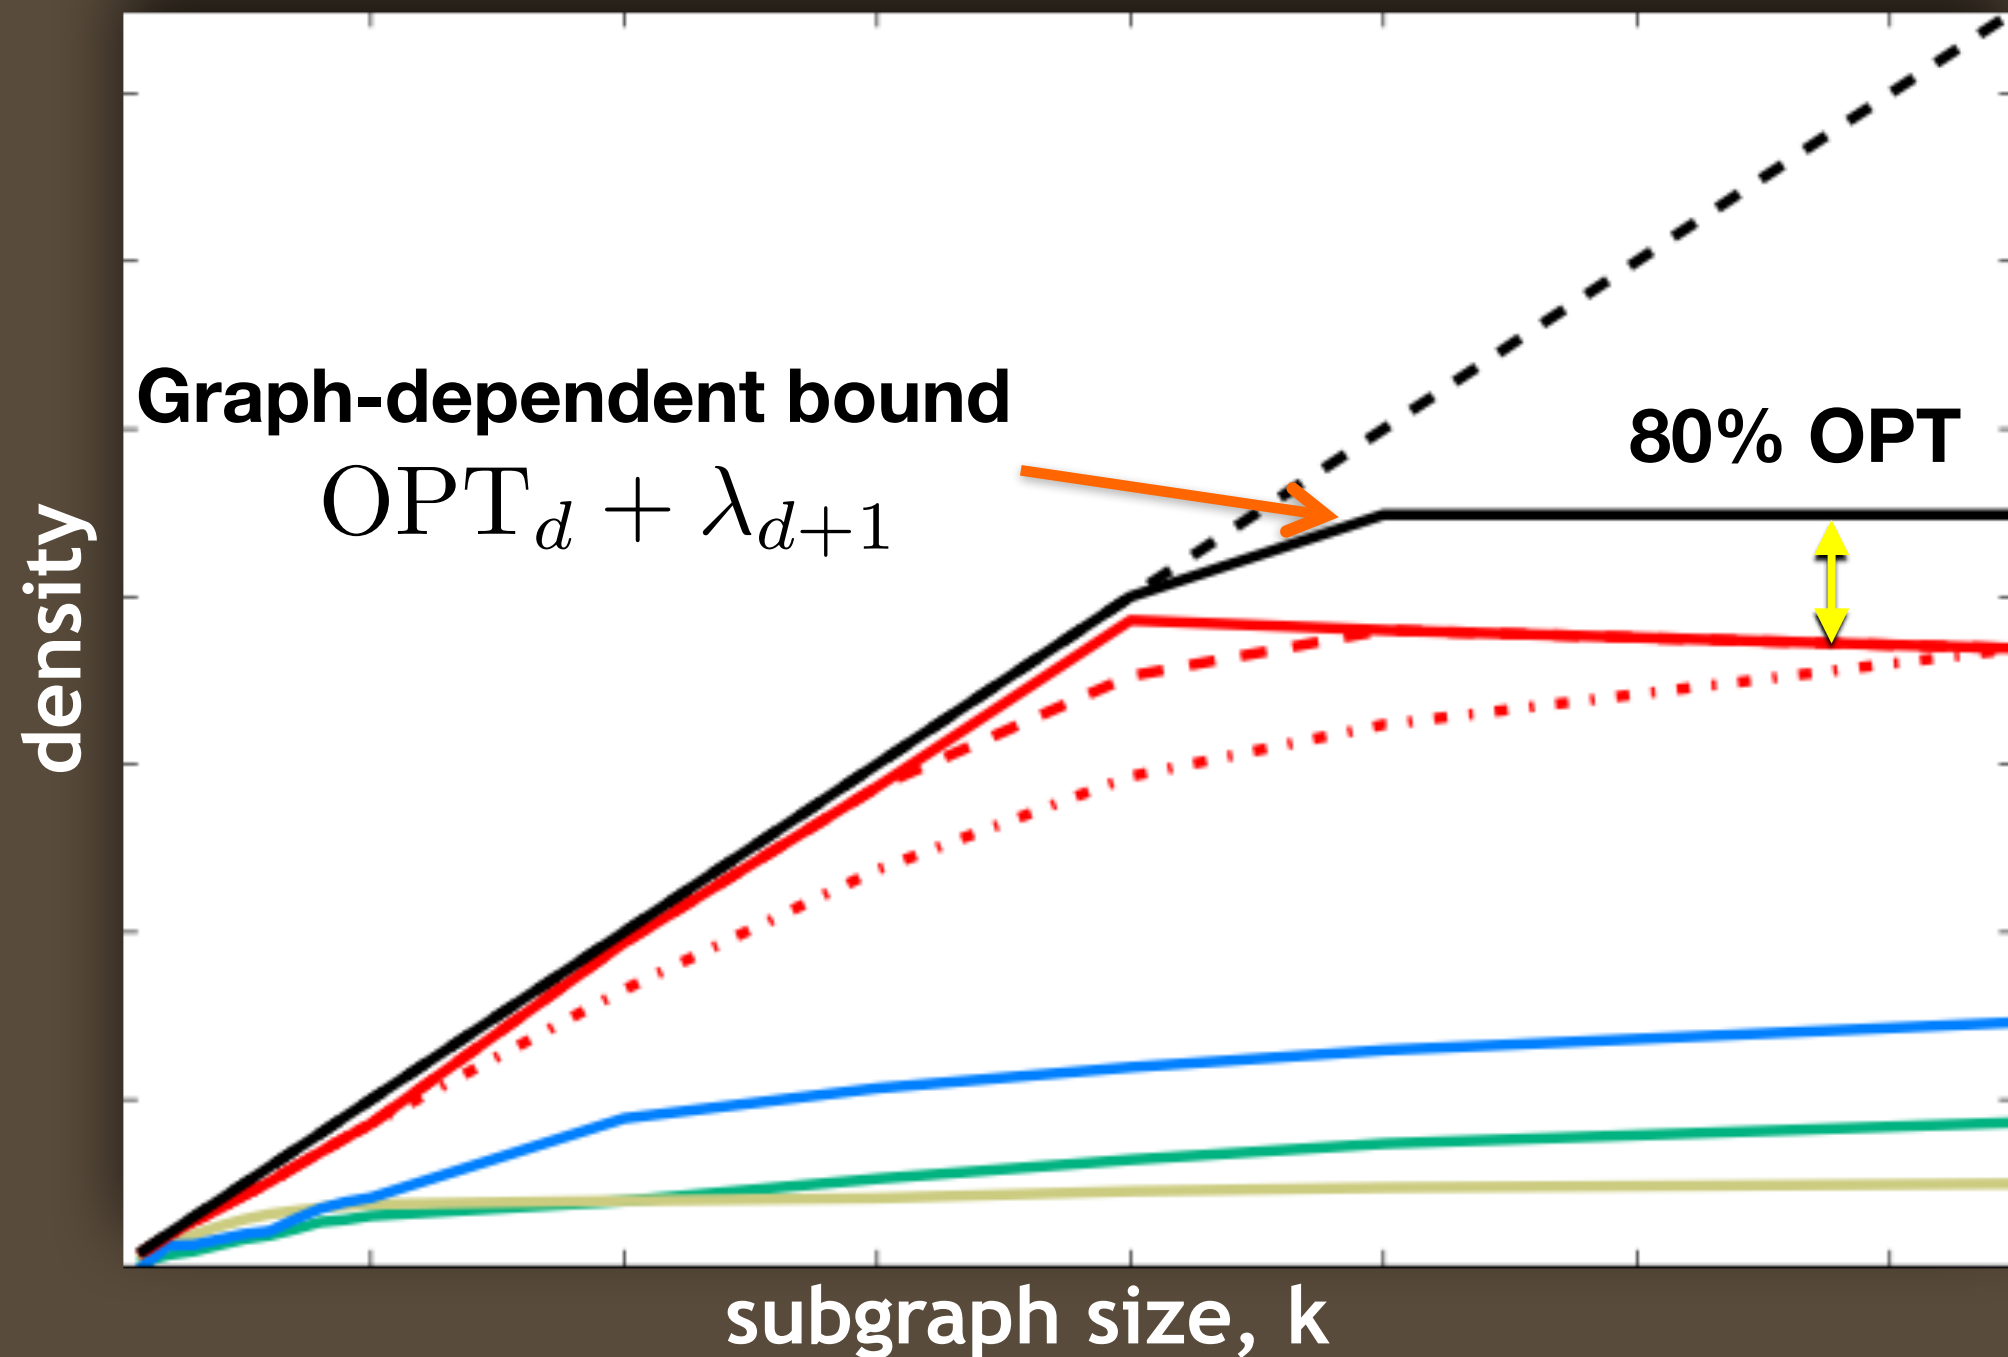
4M nodes, 35M edges



Blue: TPower JMLR'13 Green: GreedyFeige Algorithmica '01 Yellow: GreedyRavi OR'94

# com-LiveJournal graph

4M nodes, 35M edges




Blue: TPower JMLR'13   Green: GreedyFeige Algorithmica '01   Yellow: GreedyRavi OR'94




How we do it


# DkS via Quadratic Optimization



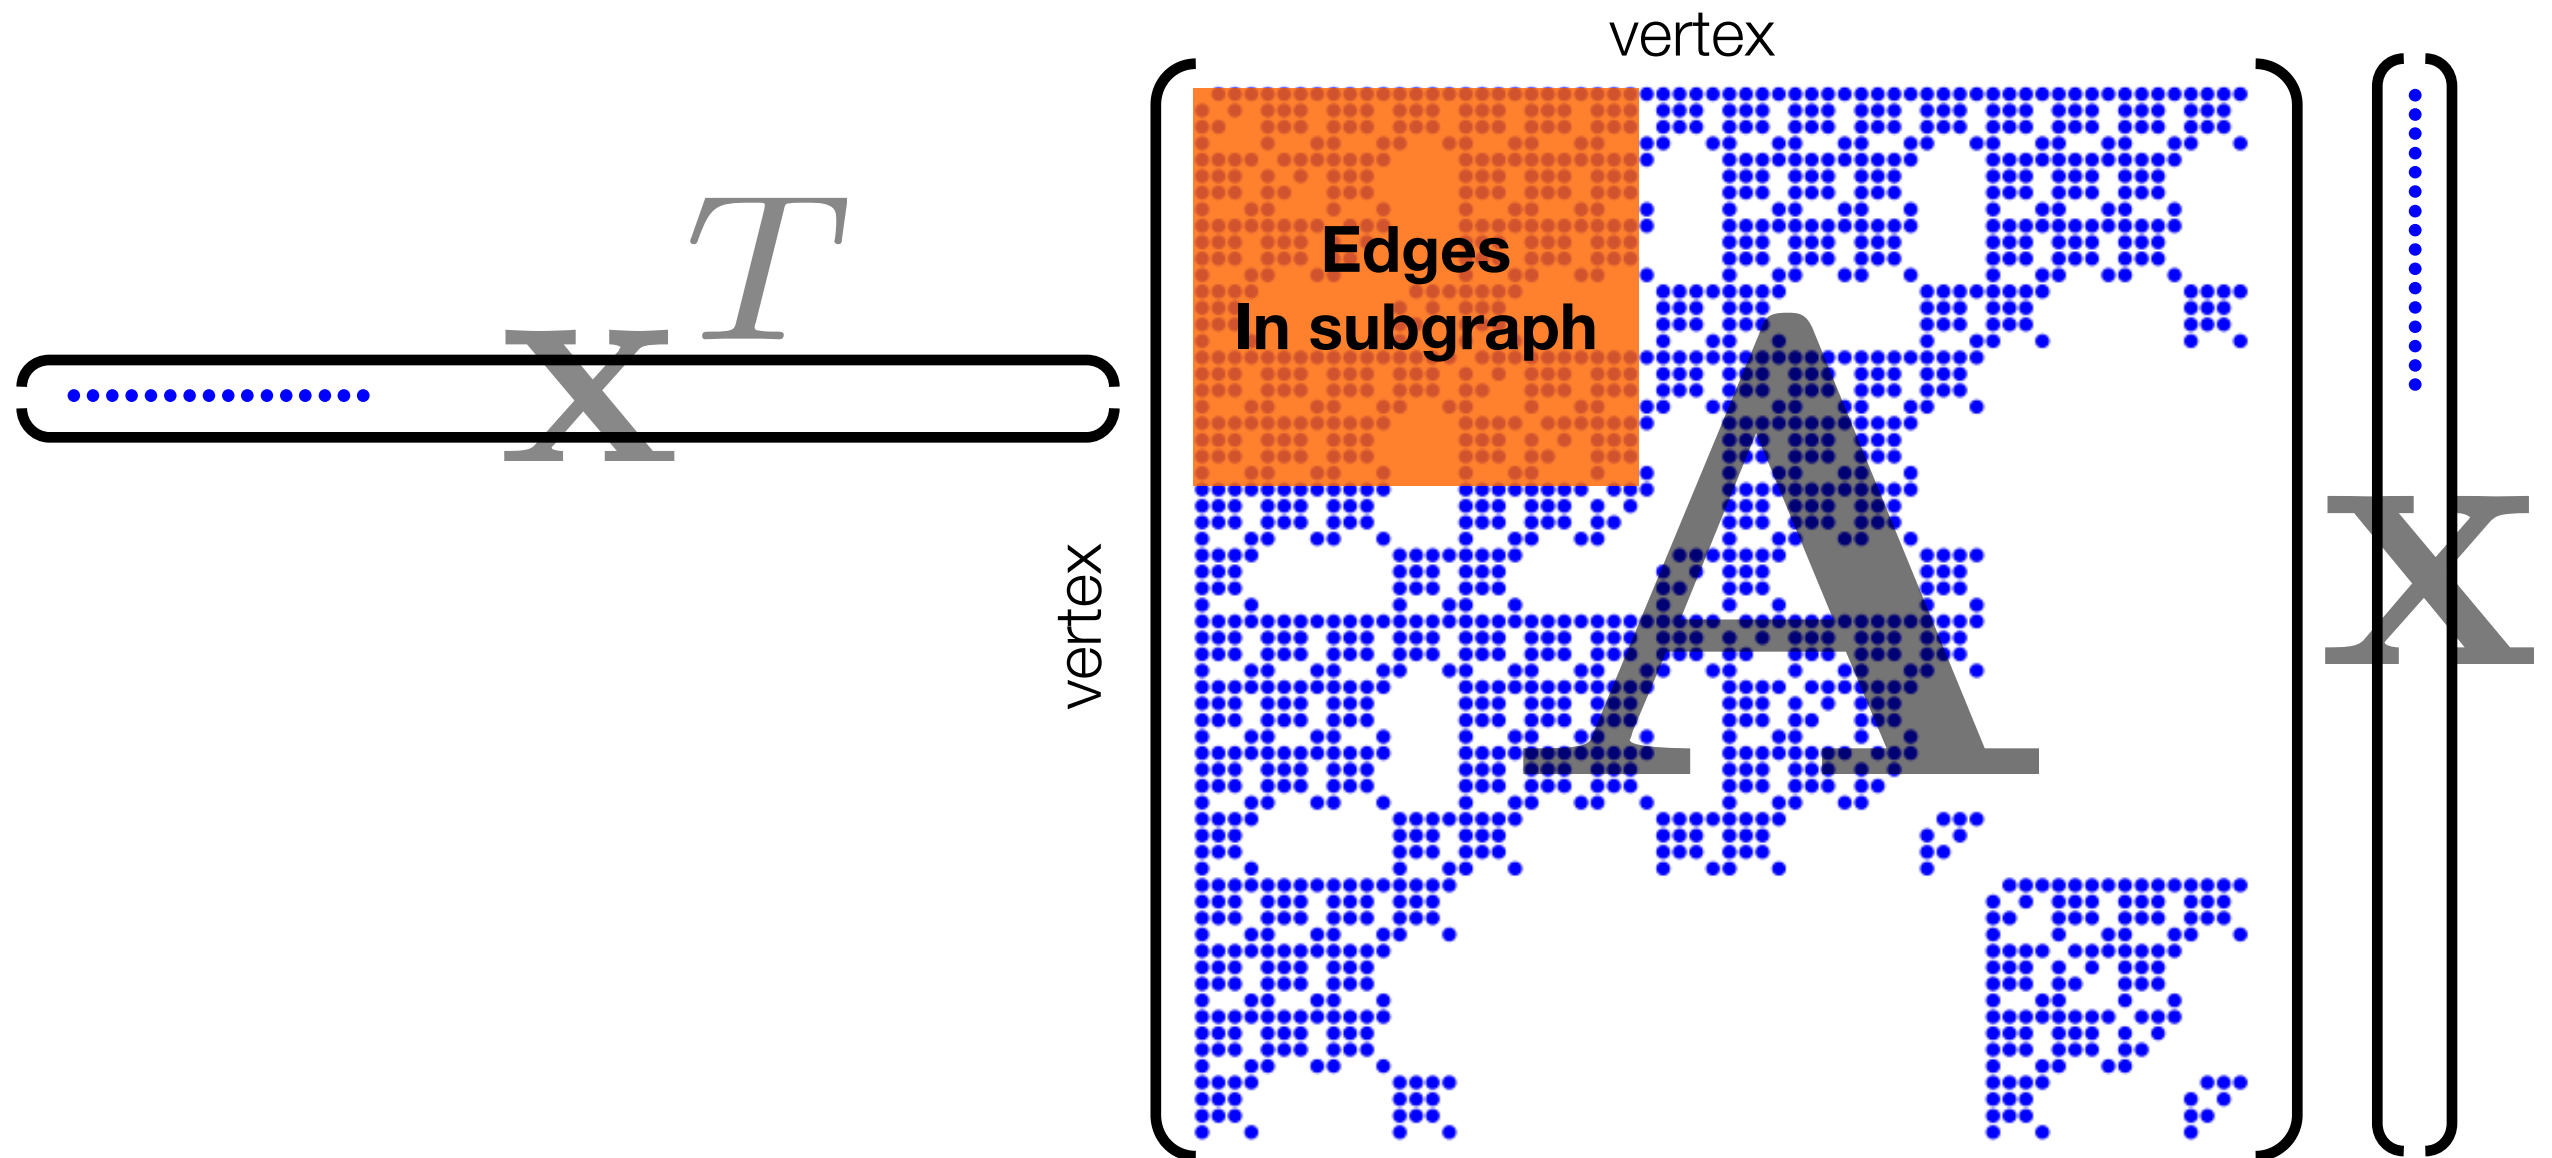
# DkS via Quadratic Optimization



# DkS via Quadratic Optimization

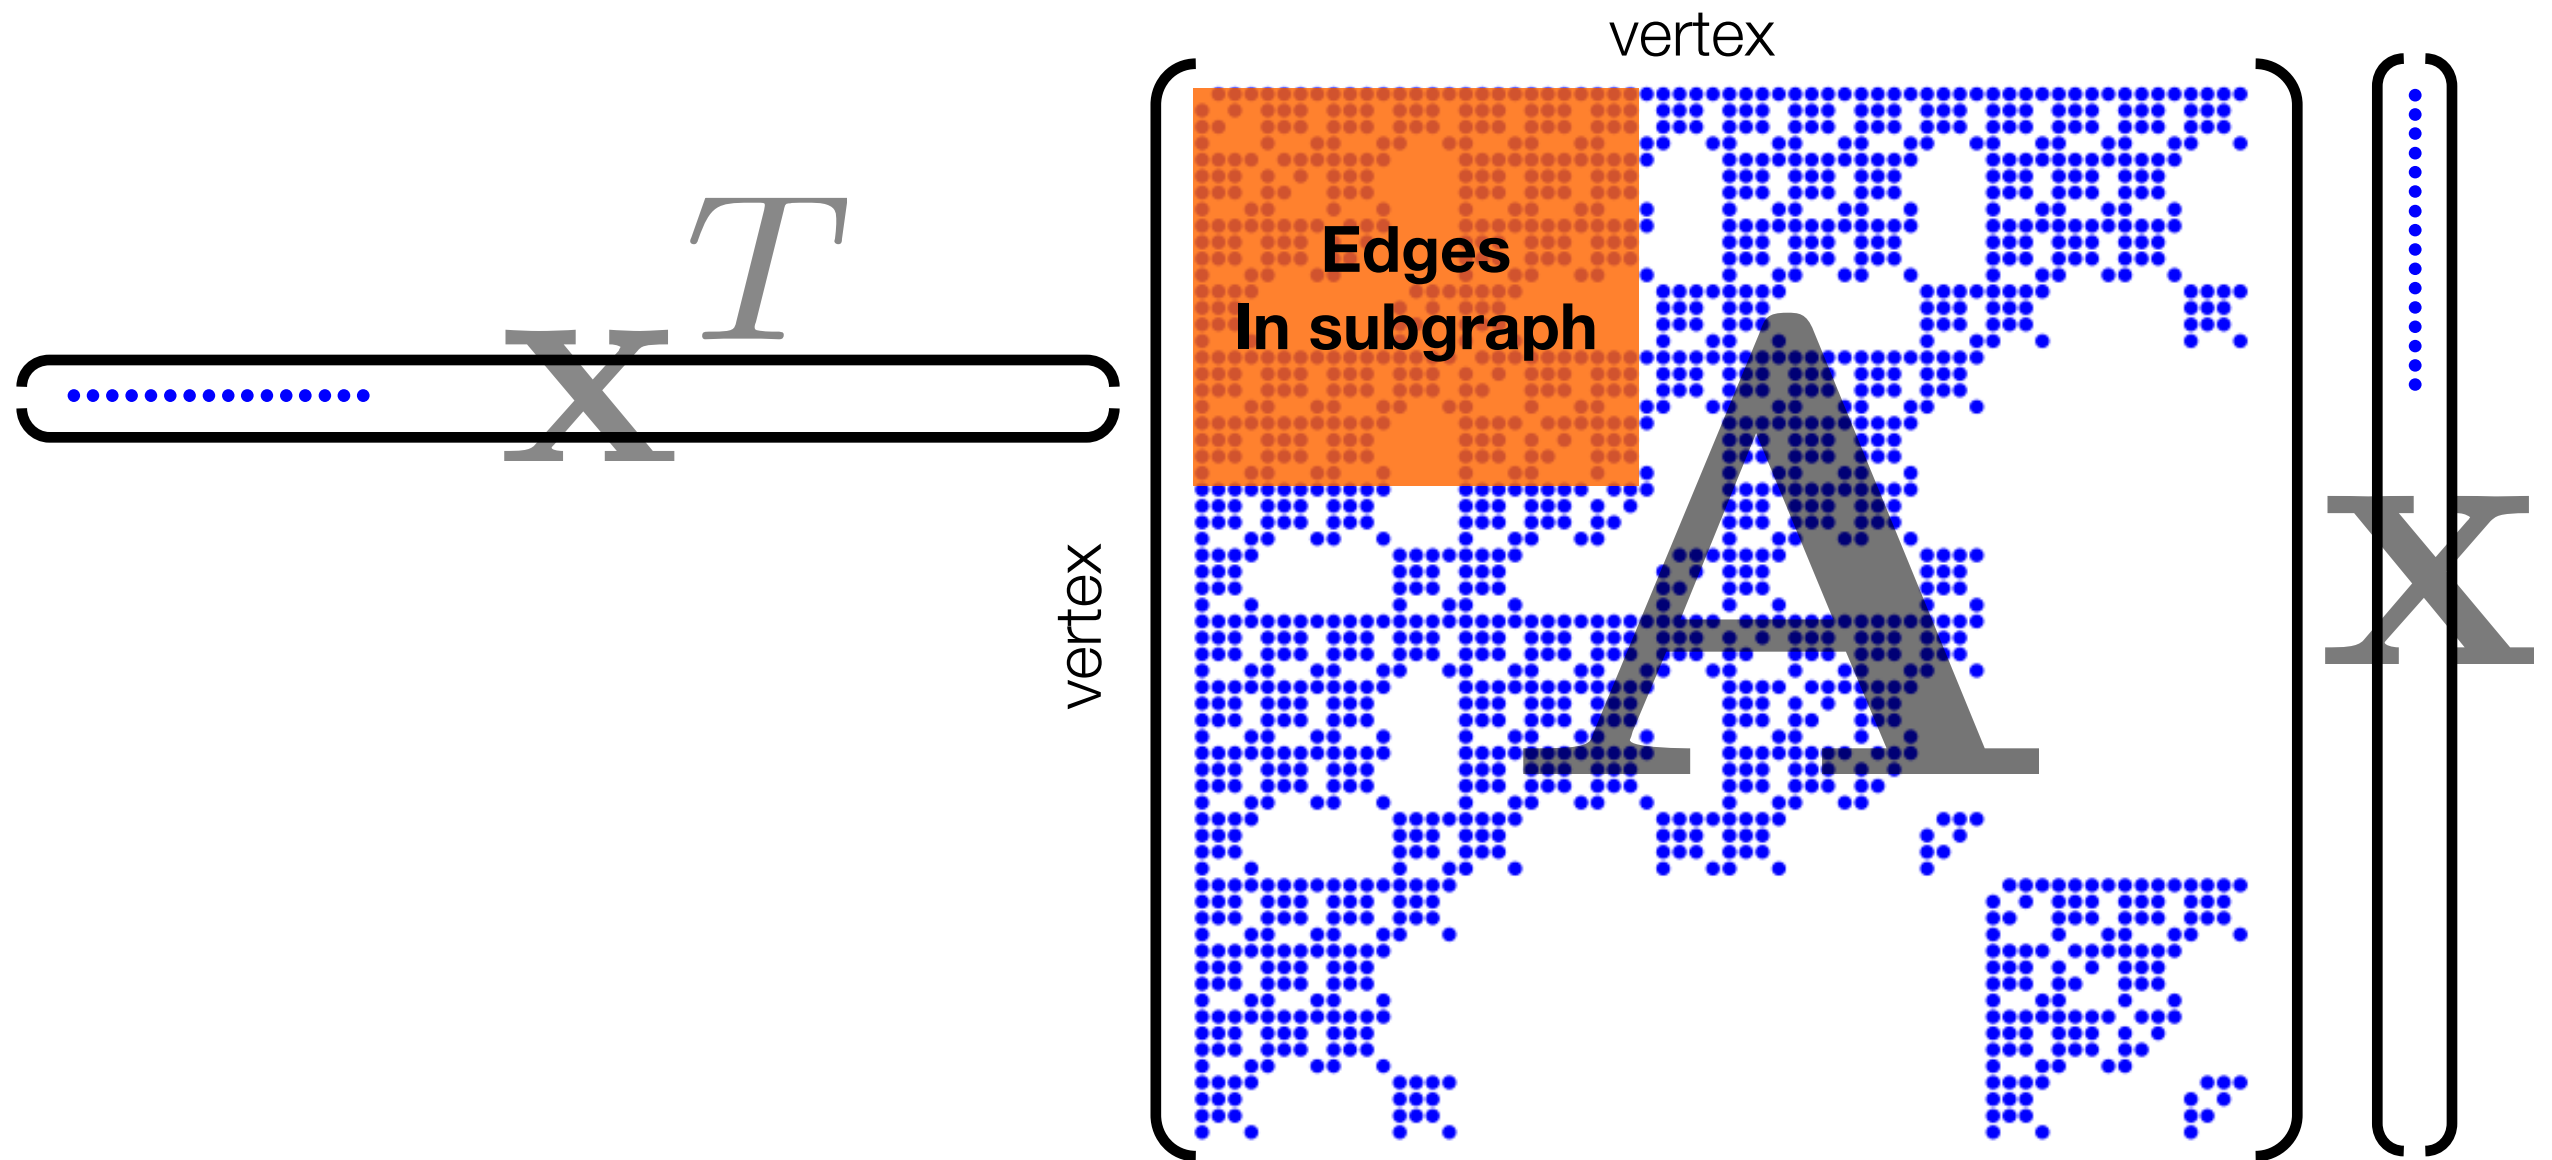


# DkS via Quadratic Optimization





# DkS via Quadratic Optimization



**DkS:** 
$$\text{OPT} = \max_{\substack{\mathbf{x} \in \{0, 1/\sqrt{k}\}^n \\ \|\mathbf{x}\|_0 = k}} \mathbf{x}^T \mathbf{A} \mathbf{x}$$

# DkS via Bilinear Optimization

**DkS:**      
$$\text{OPT} = \max_{\substack{\mathbf{x} \in \{0, 1/\sqrt{k}\}^n \\ \|\mathbf{x}\|_0 = k}} \mathbf{x}^T \mathbf{A} \mathbf{x}$$



# DkS via Bilinear Optimization

**DBkS:** 
$$\text{OPT} = \max_{\substack{\mathbf{x}, \mathbf{y} \in \{0, 1/\sqrt{k}\}^n \\ \|\mathbf{x}\|_0 = \|\mathbf{y}\|_0 = k}} \mathbf{x}^T \mathbf{A} \mathbf{y}$$

**DkS:** 
$$\text{OPT} = \max_{\substack{\mathbf{x} \in \{0, 1/\sqrt{k}\}^n \\ \|\mathbf{x}\|_0 = k}} \mathbf{x}^T \mathbf{A} \mathbf{x}$$

# DkS via Bilinear Optimization

**DBkS:** 
$$\text{OPT} = \max_{\substack{\mathbf{x}, \mathbf{y} \in \{0, 1/\sqrt{k}\}^n \\ \|\mathbf{x}\|_0 = \|\mathbf{y}\|_0 = k}} \mathbf{x}^T \mathbf{A} \mathbf{y}$$

**DkS:** 
$$\text{OPT} = \max_{\substack{\mathbf{x} \in \{0, 1/\sqrt{k}\}^n \\ \|\mathbf{x}\|_0 = k}} \mathbf{x}^T \mathbf{A} \mathbf{x}$$

# DkS via Bilinear Optimization

**DBkS:**  $\text{OPT} = \max_{\substack{\mathbf{x}, \mathbf{y} \in \{0, 1/\sqrt{k}\}^n \\ \|\mathbf{x}\|_0 = \|\mathbf{y}\|_0 = k}} \mathbf{x}^T \mathbf{A} \mathbf{y}$


**Lemma:**

**$\rho$ -approximation for DBkS =  $\frac{1}{2}\rho$ -approximation for DkS**

**DkS:**  $\text{OPT} = \max_{\substack{\mathbf{x} \in \{0, 1/\sqrt{k}\}^n \\ \|\mathbf{x}\|_0 = k}} \mathbf{x}^T \mathbf{A} \mathbf{x}$

# DkS via Bilinear Optimization

**DBkS:**  $\text{OPT} = \max_{\substack{\mathbf{x}, \mathbf{y} \in \{0, 1/\sqrt{k}\}^n \\ \|\mathbf{x}\|_0 = \|\mathbf{y}\|_0 = k}} \mathbf{x}^T \mathbf{A} \mathbf{y}$




# Low-Rank Approximation

**DBkS:**  $\text{OPT}_d = \max_{\substack{\mathbf{x}, \mathbf{y} \in \{0, 1/\sqrt{k}\}^n \\ \|\mathbf{x}\|_0 = \|\mathbf{y}\|_0 = k}} \mathbf{x}^T \mathbf{A}_d \mathbf{y}$


# Low-Rank Approximation

**DBkS:**  $\text{OPT}_d = \max_{\substack{\mathbf{x}, \mathbf{y} \in \{0, 1/\sqrt{k}\}^n \\ \|\mathbf{x}\|_0 = \|\mathbf{y}\|_0 = k}} \mathbf{x}^T \mathbf{A}_d \mathbf{y}$




# Low-Rank Approximation

**DBkS:**  $\text{OPT}_d = \max_{\substack{\mathbf{x}, \mathbf{y} \in \{0, 1/\sqrt{k}\}^n \\ \|\mathbf{x}\|_0 = \|\mathbf{y}\|_0 = k}} \mathbf{x}^T \mathbf{A}_d \mathbf{y}$



# Low-Rank Approximation

**DBkS:**  $\text{OPT}_d = \max_{\substack{\mathbf{x}, \mathbf{y} \in \{0, 1/\sqrt{k}\}^n \\ \|\mathbf{x}\|_0 = \|\mathbf{y}\|_0 = k}} \mathbf{x}^T \mathbf{A}_d \mathbf{y}$



**Efficiently solvable**



# How the Low-Rank Solver Works

**Naïvely:** Check all  $\binom{n}{k}$  subgraphs

**Rank-1 case:**  $\mathbf{A}_1 = \mathbf{v}\mathbf{v}^T$

**Q:** Maximize the product of two numbers

$$\max_{\substack{\mathbf{x}, \mathbf{y} \in \{0,1\}^n \\ \|\mathbf{x}\|_0 = \|\mathbf{y}\|_0 = k}} (\mathbf{x}^T \mathbf{v}) \cdot (\mathbf{v}^T \mathbf{y})$$

**A:** Maximize each number **individually**

# How the Rank-1 Solver Works

$$\max_{\substack{\mathbf{x}, \mathbf{y} \in \{0,1\}^n \\ \|\mathbf{x}\|_0 = \|\mathbf{y}\|_0 = k}} \left( \mathbf{x}^T \begin{array}{c} 1 \\ 2 \\ 3 \\ 4 \end{array} \right) \cdot \left( \begin{array}{c} 1 \\ 2 \\ 3 \\ 4 \end{array}^T \mathbf{y} \right)$$

**top-k set:** the k-largest coordinates of a vector,  
e.g., if  $k = 2$ , then top-2 set =  $\{3, 4\}$

**Intuition:**  $x, y$  pick the top-k set of  $v$ .

# How the Rank-2 Solver Works

$$\max(\mathbf{x}^T \begin{bmatrix} 1 & 5 \\ 2 & 2 \\ 3 & 7 \\ 4 & 0 \end{bmatrix}) (\begin{bmatrix} 1 & 5 \\ 2 & 2 \\ 3 & 7 \\ 4 & 0 \end{bmatrix}^T \mathbf{y})$$

**Intuition:**  $x, y$  pick the top- $k$  set of a vector from a 2-dimensional span.

**Q:** How many top- $k$  sets are there in a 2-dimensional span?

Based on Spannogram [Asteris, Papail., Karystinos, ISIT2011]

**Theorem:** # top- $k$  sets in a  $d$ -dimensional span:  $\binom{d}{\frac{d}{2}} \binom{n}{d} = O(n^d)$

**Spannogram:** Traverses all of them efficiently

# How the Rank-2 Solver Works

$$\max(\mathbf{x}^T \begin{bmatrix} 1 & 5 \\ 2 & 2 \\ 3 & 7 \\ 4 & 0 \end{bmatrix}) \left( \begin{bmatrix} 1 & 5 \\ 2 & 2 \\ 3 & 7 \\ 4 & 0 \end{bmatrix}^T \mathbf{y} \right)$$

**Intuition:**  $x, y$  pick the top- $k$  set of a vector from a 2-dimensional span.

Randomized algorithm

**Take random points:**  $s_1, \dots, s_{1/\epsilon^d} \in \text{span}(v_1, \dots, v_d)$

# How the Rank-2 Solver Works

$$\max(\mathbf{x}^T \begin{bmatrix} 1 & 5 \\ 2 & 2 \\ 3 & 7 \\ 4 & 0 \end{bmatrix}) \left( \begin{bmatrix} 1 & 5 \\ 2 & 2 \\ 3 & 7 \\ 4 & 0 \end{bmatrix}^T \mathbf{y} \right)$$

**Intuition:**  $x, y$  pick the top- $k$  set of a vector from a 2-dimensional span.

Randomized algorithm

**Take random points:**  $s_1, \dots, s_{1/\epsilon^d} \in \text{span}(v_1, \dots, v_d)$

Practically linear time

# Implementation

# MapReduce Implementation

```
def spannogram_mapper(self, coordinate, values):
    i = coordinate
    for j in range(int(self.options.rowcount)):
        yield j, values

def spannogram_reducer(self, coordinate, values):
    k = 5
    V = list(values)
    i = coordinate
    Vc = []
    opt_support = []
    opt_metric = 0
    for j in range(i+1, int(self.options.rowcount)):
        x = []
        Vc = []
        Vtemp = []
        # compute c_ij intersection vector
        x = [V[i][l]-V[j][l] for l in range(2)]
        # compute v_ij = Vc_ij
        Vc = [V[l][0]*x[1]-V[l][1]*x[0] for l in range(int(self.options.rowcount))]
        # find top and bottom support
        top_support_pos = zip(*heapq.nlargest(k, enumerate(Vc), key=operator.itemgetter(1)))[0]
        top_support_neg = zip(*heapq.nsmallest(k, enumerate(Vc), key=operator.itemgetter(1)))[0]
        # compute support metric
        Vtemp = [V[s] for s in top_support_pos]
        metric_pos = sum([x**2 for x in [sum(a) for a in zip(*Vtemp)]])
        Vtemp = []
        Vtemp = [V[s] for s in top_support_neg]
        metric_neg = sum([x**2 for x in [sum(a) for a in zip(*Vtemp)]])
        # find locally optimal support
        metric_list = [opt_metric, metric_pos, metric_neg]
        metric_index = metric_list.index(max(metric_list))
        opt_support = [opt_support, top_support_pos, top_support_neg][metric_index]
        opt_metric = max(metric_list)

    yield i, [opt_metric, opt_support]
```



# MapReduce Implementation

```
def spannogram_mapper(self, coordinate, values):  
    i = coordinate  
    for j in range(int(self.options.rowcount)):  
        yield j, values
```

```
def spannogram_reducer(self, coordinate, values):  
    k = 5  
    V = list(values)  
    i = coordinate
```

```
    Vc = []
```

```
    opt_support = [];
```

```
    opt_metric = 0;
```

```
    for j in range(i+1, int(self.options.rowcount)):
```

```
        x = []
```

```
        Vc = []
```

```
        Vtemp = []
```

```
        # compute c_ij intersection vector
```

```
        x = [V[i][l]-V[j][l] for l in range(2)]
```

```
        # compute v_ij = Vc_ij
```

```
        Vc = [V[l][0]*x[1]-V[l][1]*x[0] for l in range(int(self.options.rowcount))]
```

```
        # find top and bottom support
```

```
        top_support_pos = zip(*heapq.nlargest(k, enumerate(Vc), key=operator.itemgetter(1)))[0]
```

```
        top_support_neg = zip(*heapq.nsmallest(k, enumerate(Vc), key=operator.itemgetter(1)))[0]
```

```
        # compute support metric
```

```
        Vtemp = [V[s] for s in top_support_pos]
```

```
        metric_pos = sum([x**2 for x in [sum(a) for a in zip(*Vtemp)]])
```

```
        Vtemp = []
```

```
        Vtemp = [V[s] for s in top_support_neg]
```

```
        metric_neg = sum([x**2 for x in [sum(a) for a in zip(*Vtemp)]])
```

```
        # find locally optimal support
```

```
        metric_list = [opt_metric, metric_pos, metric_neg]
```

```
        metric_index = metric_list.index(max(metric_list))
```

```
        opt_support = [opt_support, top_support_pos, top_support_neg][metric_index]
```

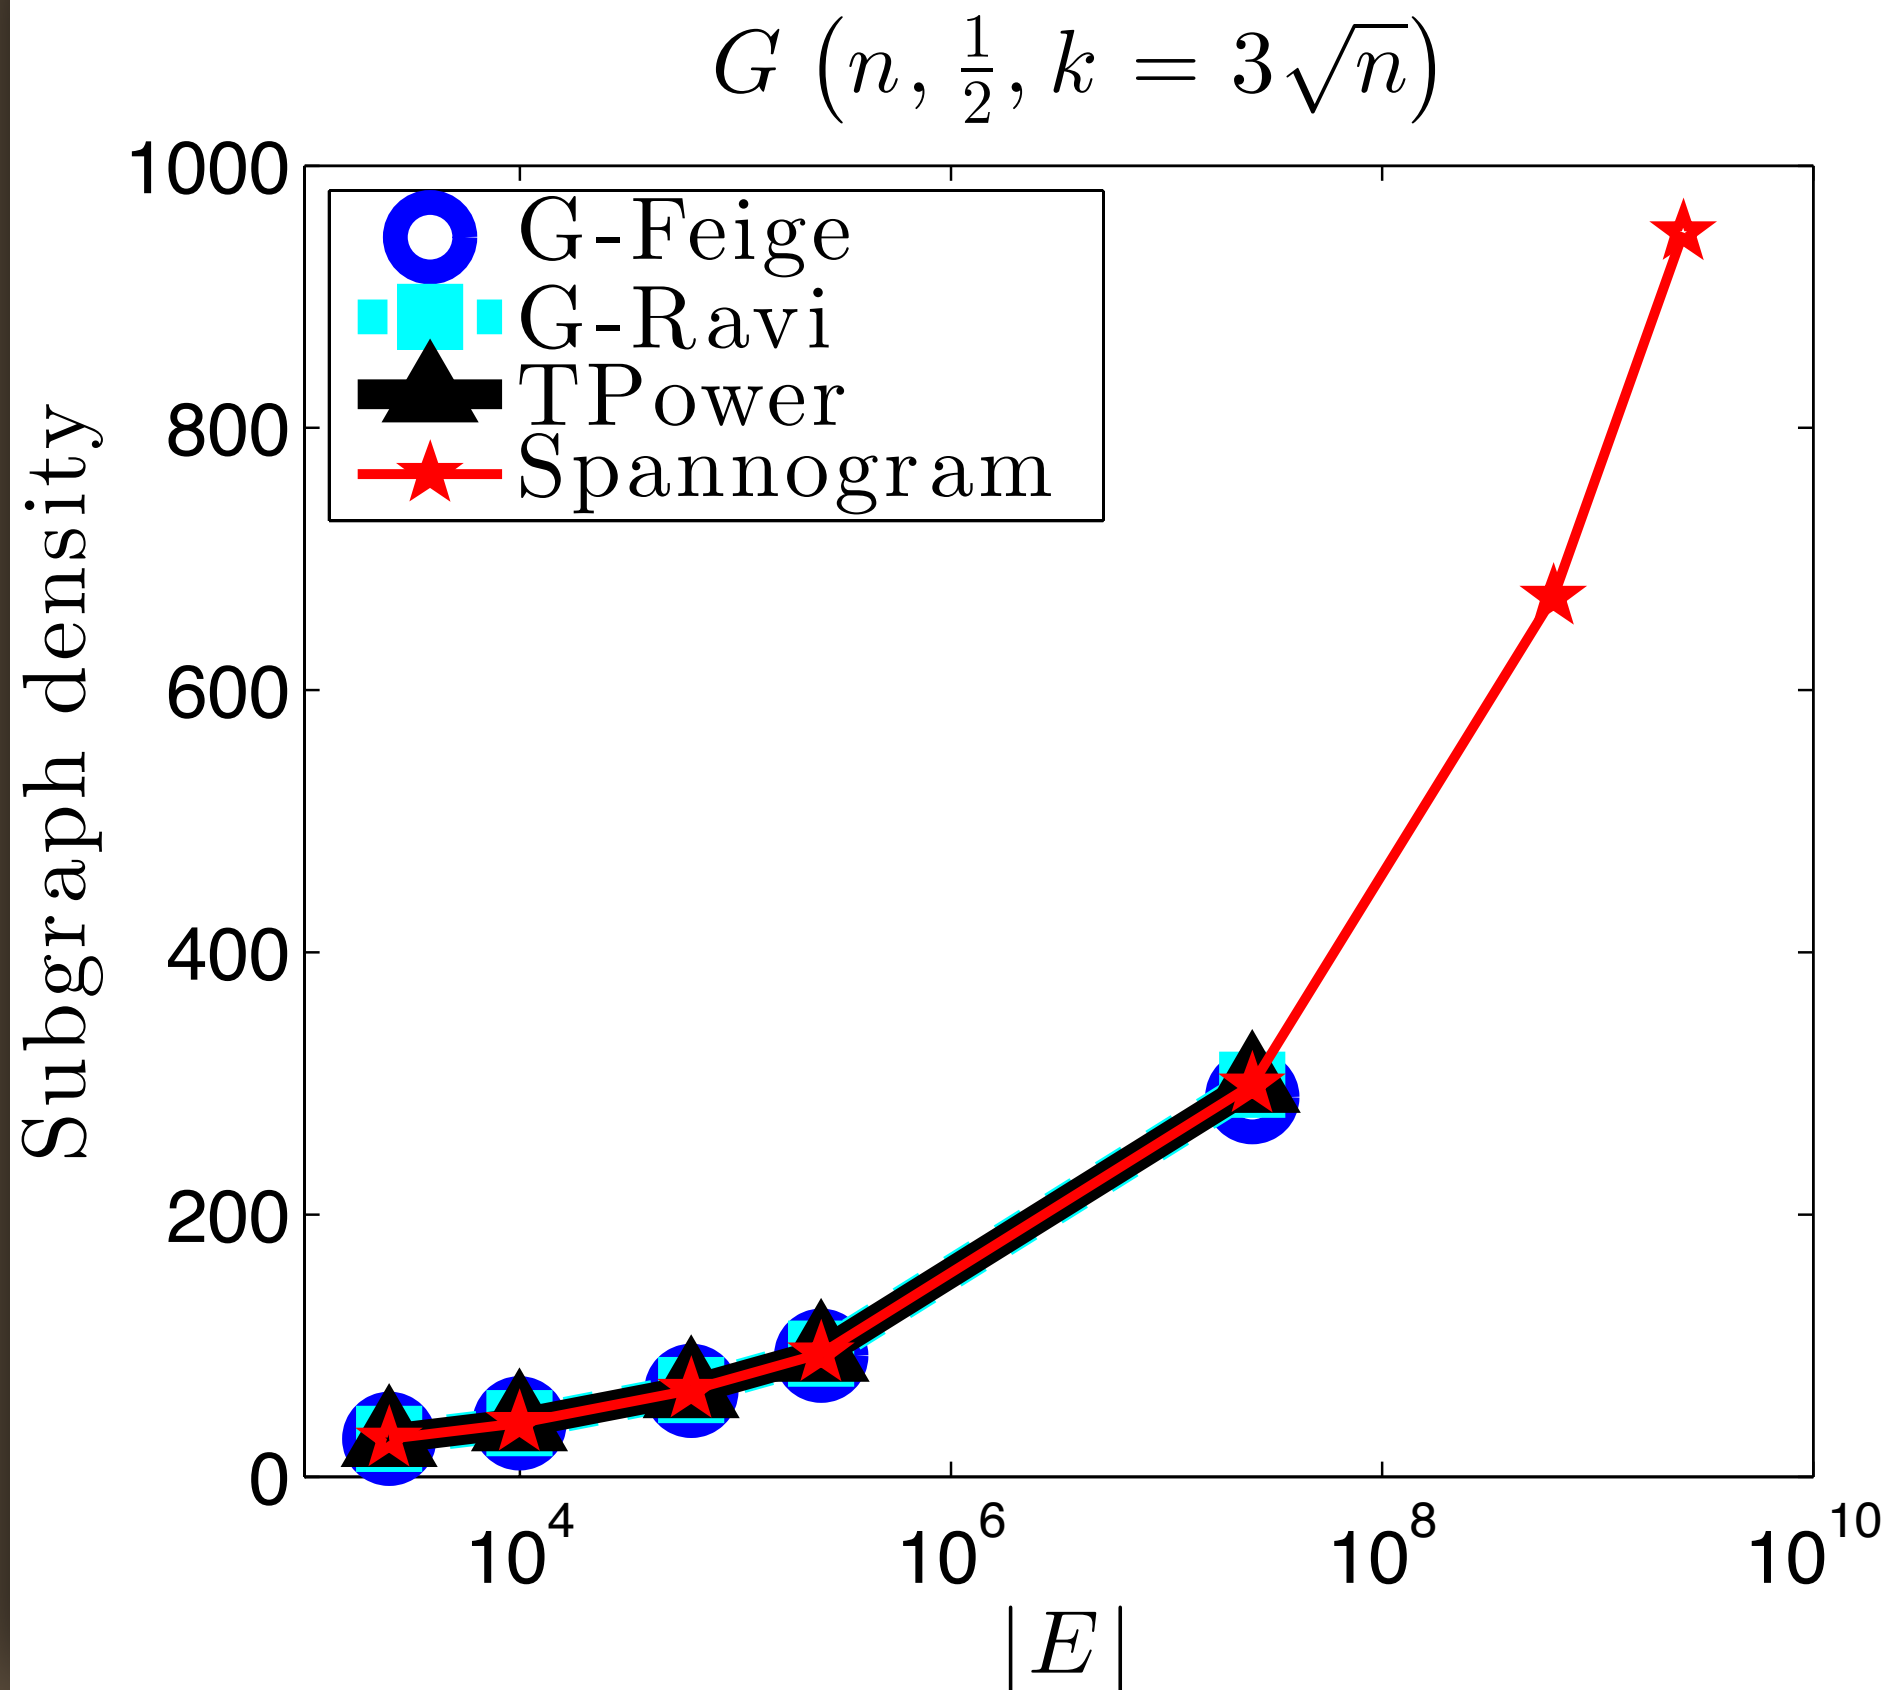
```
        opt_metric = max(metric_list)
```

```
    yield i, [opt_metric, opt_support]
```

git.io/spannogram



# Billion-scale Graphs



# Conclusions

# Conclusions

- New combinatorial approx. algorithm for DkS.

# Conclusions

- New combinatorial approx. algorithm for DkS.
- Graph-dependent spectral bounds:  
OPT within 70% in most experiments.

# Conclusions

- New combinatorial approx. algorithm for DkS.
- Graph-dependent spectral bounds:  
OPT within 70% in most experiments.
- Bound could be trivial in the worst case.

# Conclusions

- New combinatorial approx. algorithm for DkS.
- Graph-dependent spectral bounds:  
OPT within 70% in most experiments.
- Bound could be trivial in the worst case.
- Empirically outperforms previous state of the art

# Conclusions

- New combinatorial approx. algorithm for DkS.
- Graph-dependent spectral bounds:  
OPT within 70% in most experiments.
- Bound could be trivial in the worst case.
- Empirically outperforms previous state of the art

# Conclusions

- New combinatorial approx. algorithm for DkS.
- Graph-dependent spectral bounds:  
OPT within 70% in most experiments.
- Bound could be trivial in the worst case.
- Empirically outperforms previous state of the art
- Highly scalable implementation



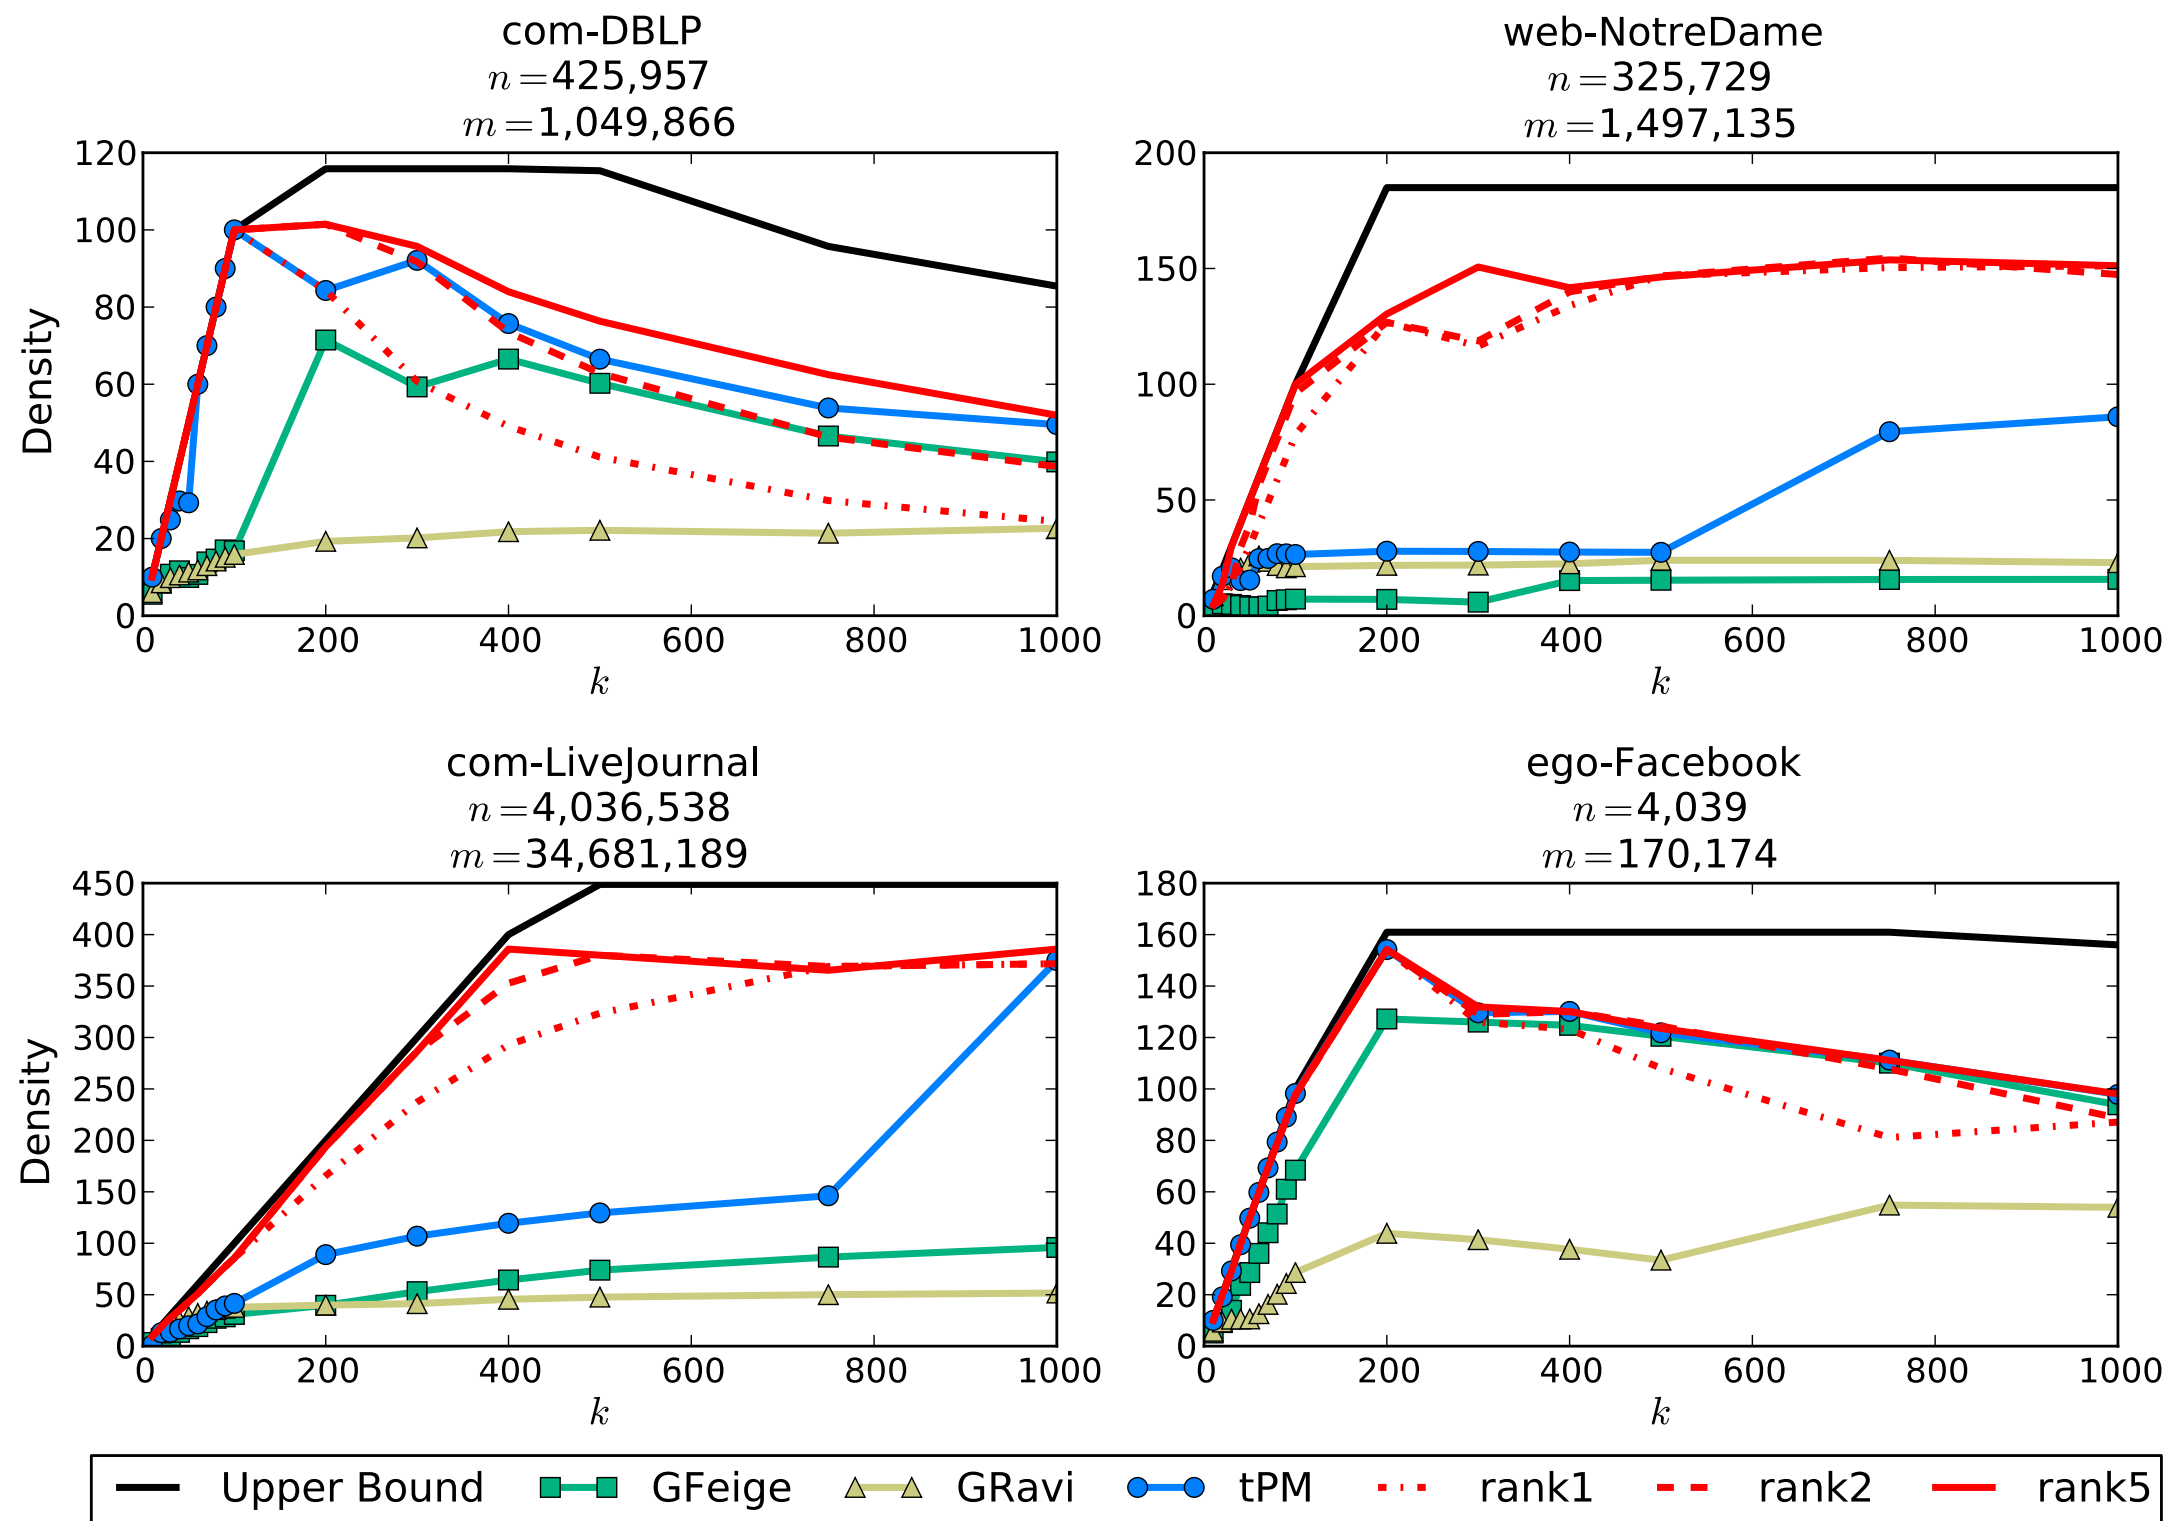
Thank you

## References

- Amazon Web Services, Elastic Map Reduce. URL <http://aws.amazon.com/elasticmapreduce/>.
- MRJob. URL <http://pythonhosted.org/mrjob/>.
- Ames, Brendan PW. *Convex relaxation for the planted clique, biclique, and clustering problems*. PhD thesis, University of Waterloo, 2011.
- Arora, Sanjeev, Karger, David, and Karpinski, Marek. Polynomial time approximation schemes for dense instances of np-hard problems. In *STOC*, 1995.
- Asahiro, Yuichi, Iwama, Kazuo, Tamaki, Hisao, and Tokuyama, Takeshi. Greedily finding a dense subgraph. *Journal of Algorithms*, 34(2):203–221, 2000.
- Asteris, Megasthenis, Papailiopoulos, Dimitris S, and Karystinos, George N. Sparse principal component of a rank-deficient matrix. In *IEEE ISIT 2011*.
- Bahmani, Bahman, Kumar, Ravi, and Vassilvitskii, Sergei. Densest subgraph in streaming and mapreduce. *Proceedings of the VLDB Endowment*, 5(5):454–465, 2012.
- Bhaskara, Aditya, Charikar, Moses, Chlamtac, Eden, Feige, Uriel, and Vijayaraghavan, Aravindan. Detecting high log-densities: an  $O(n^{1/4})$  approximation for densest k-subgraph. In *STOC*, 2010.
- Boutsidis, Christos, Mahoney, Michael W, and Drineas, Petros. An improved approximation algorithm for the column subset selection problem. In *Proceedings of the twentieth Annual ACM-SIAM Symposium on Discrete Algorithms*, pp. 968–977. Society for Industrial and Applied Mathematics, 2009.
- Cormen, Thomas H, Leiserson, Charles E, Rivest, Ronald L, and Stein, Clifford. *Introduction to algorithms*. MIT press, 2001.
- d’Aspremont, Alexandre et al. Weak recovery conditions using graph partitioning bounds. 2010.
- Dourisboure, Yon, Geraci, Filippo, and Pellegrini, Marco. Extraction and classification of dense communities in the web. In *WWW*, 2007.
- Feige, Uriel and Langberg, Michael. Approximation algorithms for maximization problems arising in graph partitioning. *Journal of Algorithms*, 41(2):174–211, 2001.
- Feige, Uriel, Peleg, David, and Kortsarz, Guy. The dense k-subgraph problem. *Algorithmica*, 29(3):410–421, 2001.
- Gibson, David, Kumar, Ravi, and Tomkins, Andrew. Discovering large dense subgraphs in massive graphs. In *PVLDB*, 2005.
- Gittens, Alex, Kambadur, Prabhanjan, and Boutsidis, Christos. Approximate spectral clustering via randomized sketching. *arXiv preprint arXiv:1311.2854*, 2013.
- Halko, Nathan, Martinsson, Per-Gunnar, and Tropp, Joel A. Finding structure with randomness: Probabilistic algorithms for constructing approximate matrix decompositions. *SIAM review*, 53(2):217–288, 2011.
- Hu, Haiyan, Yan, Xifeng, Huang, Yu, Han, Jiawei, and Zhou, Xianghong Jasmine. Mining coherent dense subgraphs across massive biological networks for functional discovery. *Bioinformatics*, 21(suppl 1):i213–i221, 2005.
- Jethava, Vinay, Martinsson, Anders, Bhattacharyya, Chiranjib, and Dubhashi, Devdatt. The lovasz theta function, svms and finding large dense subgraphs. In *NIPS*, 2012.
- Karystinos, George N and Liavas, Athanasios P. Efficient computation of the binary vector that maximizes a rank-deficient quadratic form. *IEEE Trans. IT*, 56(7):3581–3593, 2010.
- Khot, Subhash. Ruling out ptas for graph min-bisection, densest subgraph and bipartite clique. In *FOCS*, 2004.
- Lin, Jimmy and Schatz, Michael. Design patterns for efficient graph algorithms in mapreduce. In *Proceedings of the Eighth Workshop on Mining and Learning with Graphs*, pp. 78–85. ACM, 2010.
- Mahoney, Michael W and Drineas, Petros. Cur matrix decompositions for improved data analysis. *Proceedings of the National Academy of Sciences*, 106(3):697–702, 2009.
- Meng, Xiangrui and Mahoney, Michael W. Robust regression on mapreduce. *ICML 2013, (to appear)*.
- Miller, B, Bliss, N, and Wolfe, P. Subgraph detection using eigenvector l1 norms. In *NIPS*, 2010.
- Papailiopoulos, Dimitris S, Dimakis, Alexandros G, and Korokythakis, Stavros. Sparse pca through low-rank approximations. *arXiv preprint arXiv:1303.0551*, 2013.
- Ravi, Sekharipuram S, Rosenkrantz, Daniel J, and Tayi, Giri K. Heuristic and special case algorithms for dispersion problems. *Operations Research*, 42(2):299–310, 1994.
- Rokhlin, Vladimir, Szlam, Arthur, and Tygert, Mark. A randomized algorithm for principal component analysis. *SIAM Journal on Matrix Analysis and Applications*, 31(3):1100–1124, 2009.
- Saha, Barna, Hoch, Allison, Khuller, Samir, Raschid, Louiqa, and Zhang, Xiao-Ning. Dense subgraphs with restrictions and applications to gene annotation graphs. In *Research in Computational Molecular Biology*, pp. 456–472. Springer, 2010.
- Srivastav, Anand and Wolf, Katja. *Finding dense subgraphs with semidefinite programming*. Springer, 1998.
- Suzuki, Akiko and Tokuyama, Takeshi. Dense subgraph problems with output-density conditions. In *Algorithms and Computation*, pp. 266–276. Springer, 2005.
- Wyner, Aaron D. Random packings and coverings of the unit n-sphere. *Bell System Technical Journal*, 46(9):2111–2118, 1967.
- Yuan, Xiao-Tong and Zhang, Tong. Truncated power method for sparse eigenvalue problems. *arXiv preprint arXiv:1112.2679*, 2011.

Backup slides

# Other experiments



# Randomized Algorithm

## Step 1

**Take random points:**  $s_1, \dots, s_{1/\epsilon^d} \in \text{span}(v_1, \dots, v_d)$

## Step 2

**Find largest k entries:**  $\text{top}_k(s_i)$

## Step 3

**Compute density of corresponding subgraph**

# Randomized Algorithm

## Step 1

**Take random points:**  $s_1, \dots, s_{1/\epsilon^d} \in \text{span}(v_1, \dots, v_d)$

## Step 2

**Find largest k entries:**  $\text{top}_k(s_i)$

## Step 3

**Compute density of corresponding subgraph**

Practically linear time